


gemm vector layout W dequant KV dequant

relative energy/token

llama2 Prefill


batch=1



relative energy/token

llama3 Prefill

batch=1




relative energy/token

llama2 Decode

batch=1

batch=4

batch=64




relative energy/token

llama3 Decode

batch=1

batch=4

batch=64



context length